

ESE 6510

## Generative Models

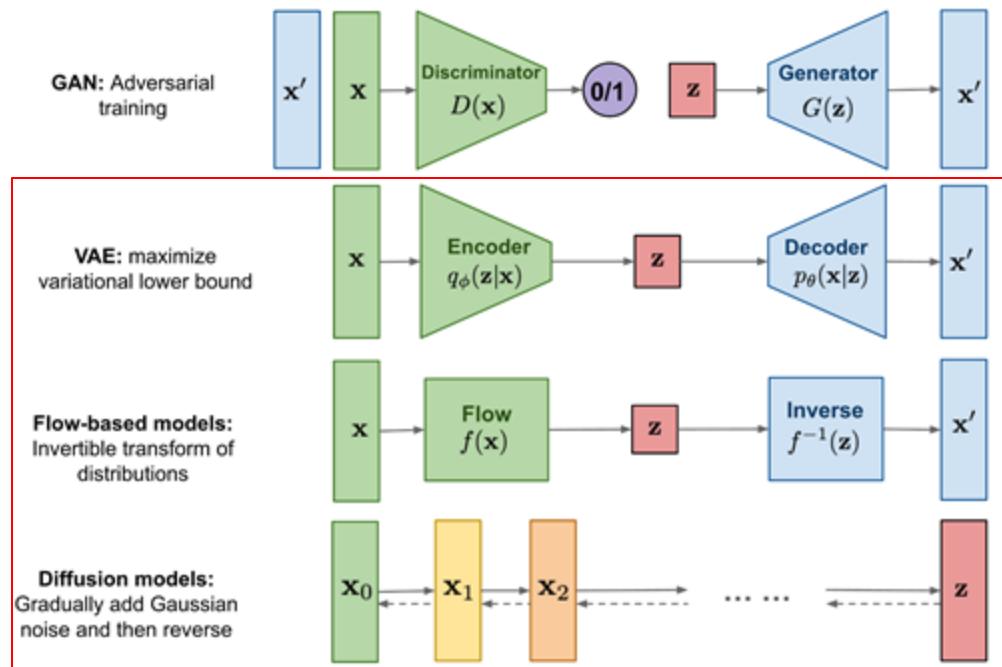
Presented by: Chunwei Xing



Vincent van Gogh, 1889

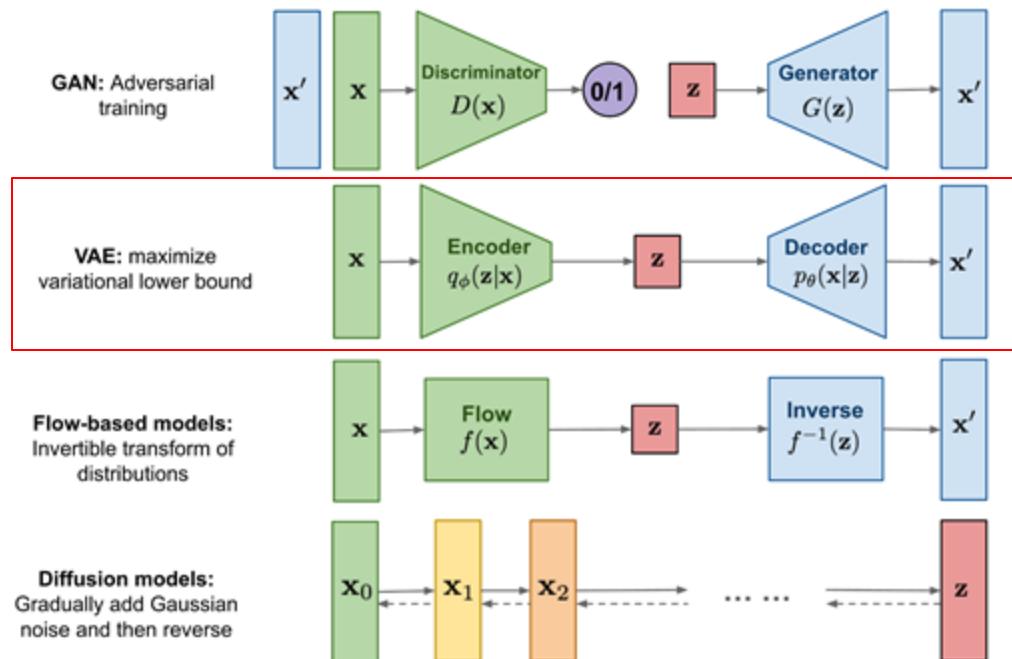
# Generative Models

- Variational Autoencoder
- Diffusion Models
- Flow-based Models



# Generative Models

- Variational Autoencoder
- Diffusion Models
- Flow-based Models



# Autoencoder - A Brief History

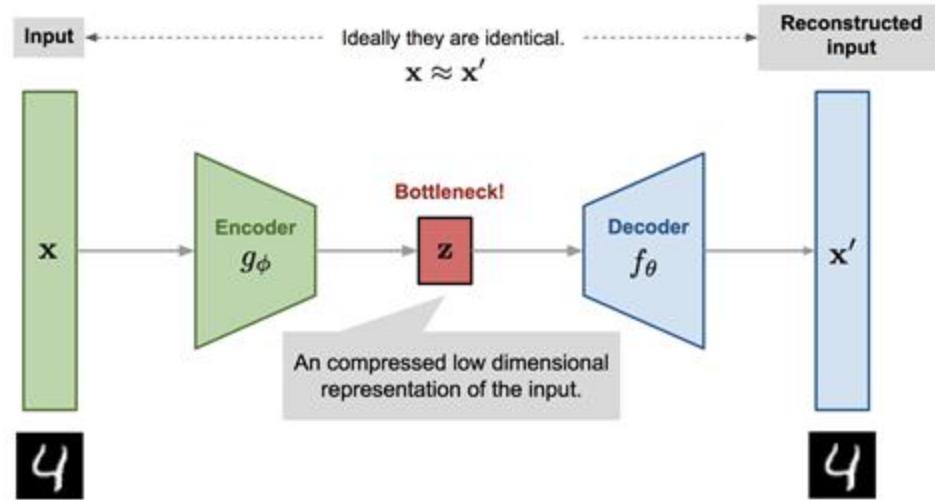
- 1982 — PCA: Oja showed PCA is equivalent to a 1-hidden-layer linear neural net
- 1989–1991 — Nonlinear PCA: Baldi & Hornik (1989) and Kramer (1991) generalized PCA to neural “autoassociative” networks
- Mid-late 1980s — Auto-association: The idea to run a neural net in “auto-association mode” (1986) was implemented for speech (1987–88) and images (1987).
- Early 1990s — Applications: dimensionality reduction/feature learning
- 2006 — Deep revival via pretraining: Hinton & Salakhutdinov popularized deep autoencoders using layer-wise pretraining
- Nowadays: generative modeling using VAE for large-scale generative AI

# Autoencoders

- ❑ Compressed representation
- ❑ Unsupervised learning
- ❑ Encoder network:  $\mathbf{z} = g_\phi(\mathbf{x})$
- ❑ Decoder network:  $\mathbf{x}' = f_\theta(g_\phi(\mathbf{x}))$
- ❑ Reconstruction loss:

$$L_{AE}(\theta, \phi) = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}^{(i)} - f_\theta(g_\phi(\mathbf{x}^{(i)})))^2$$

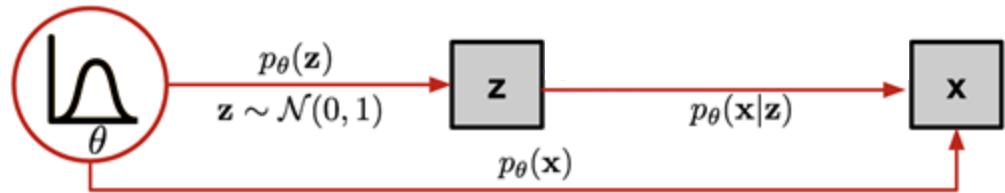
- ❑ But it's not generative!



# Variational Autoencoder

- Prior  $p_\theta(z)$
- Likelihood  $p_\theta(x|z)$
- Maximize the log-likelihood:

$$\theta^* = \arg \max_{\theta} \sum_{i=1}^n \log p_{\theta}(\mathbf{x}^{(i)})$$



- Compute  $\log p_{\theta}(x^{(i)}) = \log \int p_{\theta}(x^{(i)} | z)p(z)dz$
- What's the issue?
  - No closed-form expression for general neural network parameterizations
  - Expensive to approximate the integral over many latents for each data point

# Variational Inference

- **Theorem:** the likelihood can be written as  $\log p_\theta(x) = \max_{q(\cdot|x): q(\cdot|x) \geq 0, \int q(z|x)dz=1} \int q(z|x) \log \frac{p_\theta(x, z)}{q(z|x)} dz.$  and the maximizing distribution is given by  $q^*(z|x) = p_\theta(z|x)$

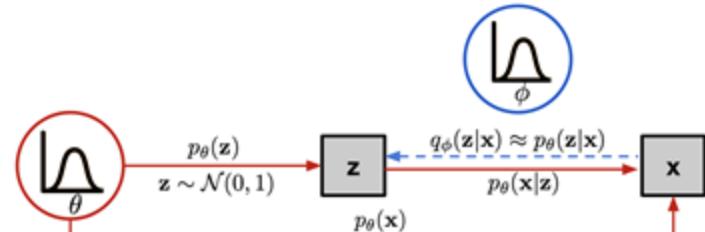
- Therefore, the new objective is given by

$$\max_{\theta} \sum_{i=1}^n \log p_\theta(x^{(i)}) = \max_{\theta} \max_{q(\cdot|x^{(i)}), \forall i} \sum_{i=1}^n \int q(z|x^{(i)}) \log \frac{p_\theta(x^{(i)}, z)}{q(z|x^{(i)})} dz$$

- Approximate the posterior with neural networks parameterized by  $q_\phi(z|x)$

$$\max_{\theta} \max_{\phi} \sum_{i=1}^n \int q_\phi(z|x^{(i)}) \log \frac{p_\theta(x^{(i)}, z)}{q_\phi(z|x^{(i)})} dz$$

- Is the new objective tractable now?



# Proof - VI Theorem

$$\begin{aligned}\log p_\theta(x) &= \int q(z|x) \log p_\theta(x) dz \\&= \int q(z|x) \log \frac{p_\theta(x, z)}{p_\theta(z|x)} dz \\&= \int q(z|x) \log \frac{p_\theta(x, z)}{q(z|x)} \cdot \frac{q(z|x)}{p_\theta(z|x)} dz \\&= \underbrace{\int q(z|x) \log \frac{p_\theta(x, z)}{q(z|x)} dz}_{\text{Evidence Lower Bound (ELBO)}} + \underbrace{\int q(z|x) \log \frac{q(z|x)}{p_\theta(z|x)} dz}_{\text{KL}[q(z|x)\|p_\theta(z|x)]}\end{aligned}$$

□ Given that  $\int q(z|x) dz = 1$

$$\begin{aligned}\max_{q(\cdot|x)} \int q(z|x) \log \frac{p_\theta(x, z)}{q(z|x)} dz &= \max_{q(\cdot|x)} \log p_\theta(x) - \text{KL}[q(z|x)\|p_\theta(z|x)] \\&= \log p_\theta(x) - \min_{q(\cdot|x)} \text{KL}[q(z|x)\|p_\theta(z|x)] \\&= \log p_\theta(x)\end{aligned}$$

□ Given that KL divergence non-negative and

$$\text{KL}(q\|p) = 0 \text{ iff } p = q$$

# Variational Autoencoder

$$\mathcal{L}(\theta, \phi; \mathbf{x}) = \underbrace{\mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\log p_\theta(\mathbf{x} | \mathbf{z})]}_{\text{reconstruction term}} - \underbrace{\text{KL}(q_\phi(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z}))}_{\text{regularization to prior}}$$

- Learning objective: maximize the ELBO
  - Maximize the likelihood of generating real data (decoder)
  - Minimize the difference between the prior and posterior distributions (encoder)
- An example
  - Encoder:  $q_\phi(\mathbf{z} | \mathbf{x}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}_\phi(\mathbf{x}), \text{diag}(\boldsymbol{\sigma}_\phi^2(\mathbf{x})))$ .
  - Prior:  $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$ .
  - Decoder:  $p_\theta(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_\theta(\mathbf{z}), \eta \mathbf{I})$ .
  - Reconstruction term:  $\mathbb{E}_{q_\phi} [\log p_\theta(\mathbf{x} | \mathbf{z})] \approx -\frac{1}{2\eta M} \sum_{m=1}^M \|\mathbf{x} - \boldsymbol{\mu}_\theta(\mathbf{z}^{(m)})\|_2^2 + \text{const.}$
  - Regularization term:  $\text{KL}(\mathcal{N}(\boldsymbol{\mu}, \text{diag } \boldsymbol{\sigma}^2) \| \mathcal{N}(\mathbf{0}, \mathbf{I})) = \frac{1}{2} \sum_{j=1}^d (\sigma_j^2 + \mu_j^2 - 1 - \log \sigma_j^2)$ .

# VAE - Reparameterization Tricks

- We can estimate gradients wrt.  $\theta$  using MC estimation

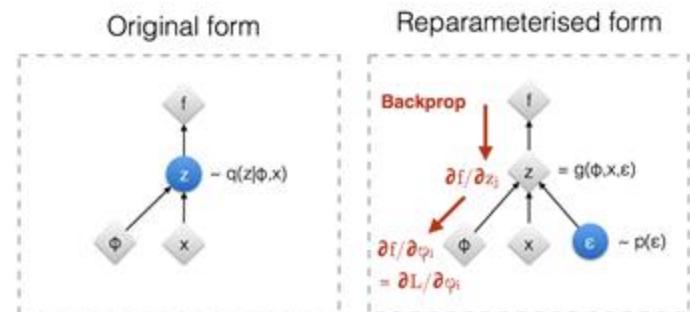
$$\nabla_{\theta} \mathcal{L}(\theta, \phi; \mathbf{x}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\nabla_{\theta} \log p_{\theta}(\mathbf{x} | \mathbf{z})] \approx \frac{1}{M} \sum_{m=1}^M \nabla_{\theta} \log p_{\theta}(\mathbf{x} | \mathbf{z}^{(m)}).$$

- But not wrt.  $\phi$

$$\nabla_{\phi} \mathbb{E}_{z \sim q_{\phi}(z|x)} [\log p_{\theta}(x, z) - \log q_{\phi}(z|x)] \neq \mathbb{E}_{z \sim q_{\phi}(z|x)} \nabla_{\phi} [\log p_{\theta}(x, z) - \log q_{\phi}(z|x)]$$

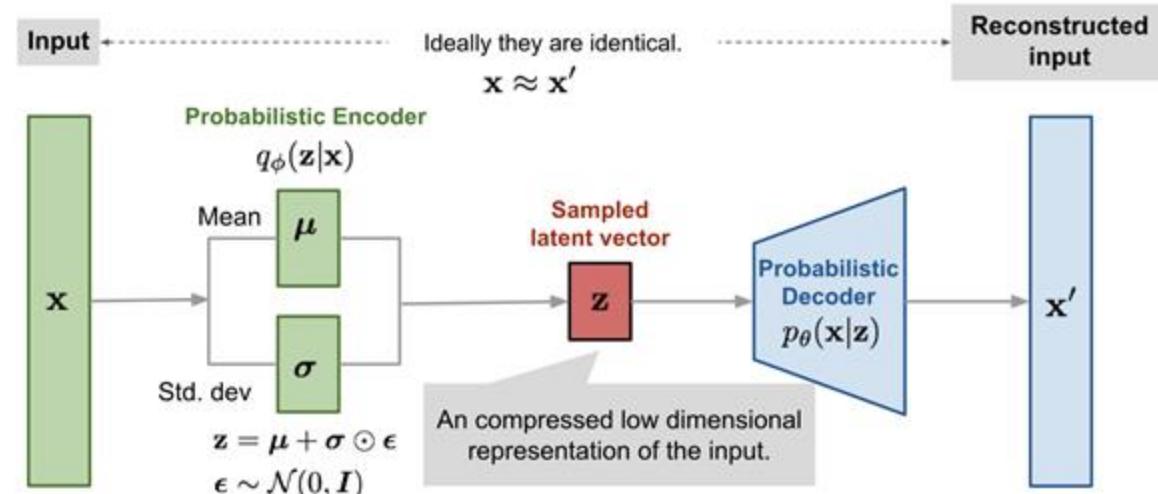
- Reparameterization and sample  $z_{\phi} = g(\epsilon, \phi, x) = \mu_{\phi}(x) + \sigma_{\phi}(x)\epsilon$  for  $\epsilon \sim \mathcal{N}(0, I)$
- Then we can estimate the gradients

$$\begin{aligned}\nabla_{\phi} \mathcal{L}_{\theta, \phi}(x) &= \nabla_{\phi} \mathbb{E}_{z \sim q_{\phi}(z|x)} [\log p_{\theta}(x, z) - \log q_{\phi}(z|x)] \\ &= \nabla_{\phi} \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [\log p_{\theta}(x, z_{\phi}) - \log q_{\phi}(z_{\phi}|x)]\end{aligned}$$



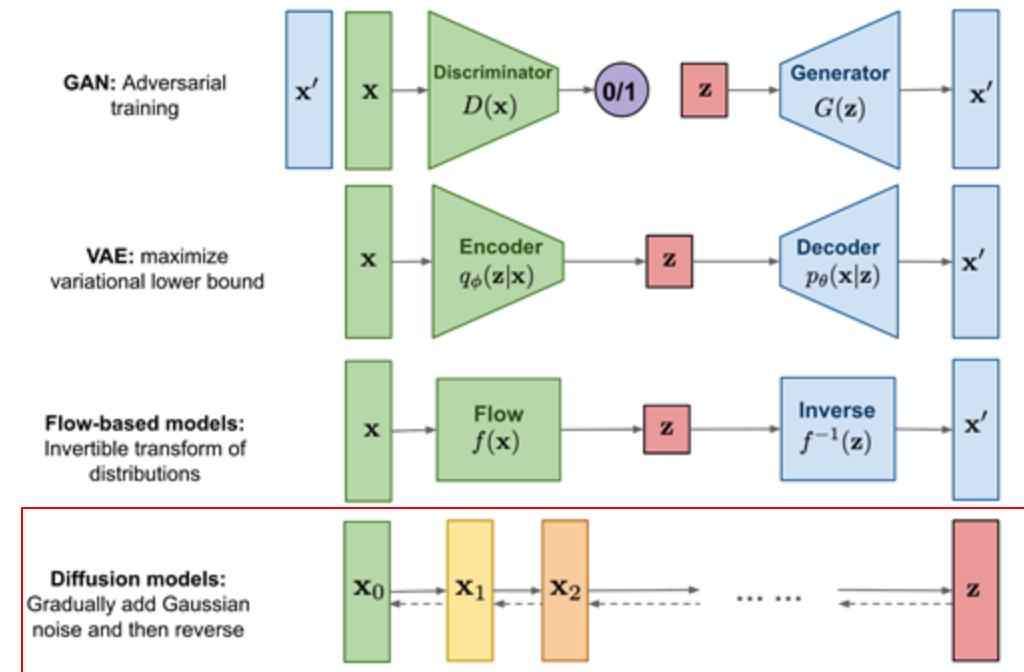
# Variational Autoencoder

- Beta-VAE
- Joint-VAE
- VQ-VAE
- ...

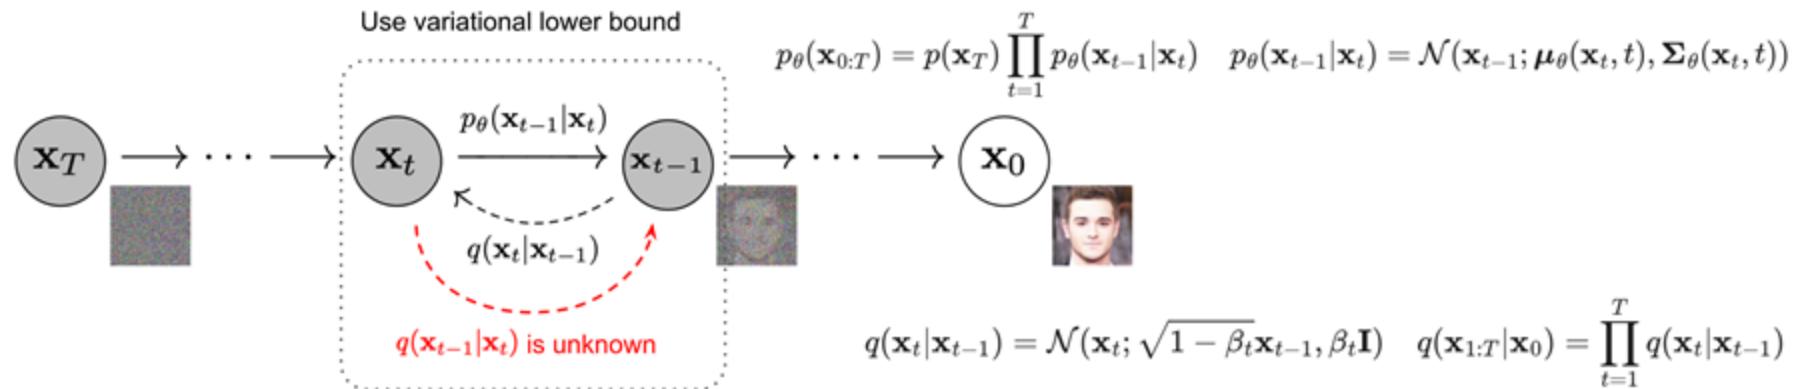


# Generative Models

- ❑ Variational Autoencoder
- ❑ Diffusion Models
- ❑ Flow-based Models



# Diffusion Models



- ❑ Forward diffusion process: define a **Markov chain** of diffusion steps to slowly add random noise to data
- ❑ Reverse diffusion process: construct desired data samples from the noise
- ❑ Connections to the VAE? Encoder? Decoder? Latents? Prior? Posterior?

# Forward Diffusion Process

- Compute  $q(x_t|x_0)$
- Given:  $q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$      $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$

$$\begin{aligned}\mathbf{x}_t &= \sqrt{\alpha_t}\mathbf{x}_{t-1} + \sqrt{1 - \alpha_t}\boldsymbol{\epsilon}_{t-1} \quad \text{reparameterization trick: } \boldsymbol{\epsilon}_{t-1}, \boldsymbol{\epsilon}_{t-2}, \dots \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \\ &= \sqrt{\alpha_t\alpha_{t-1}}\mathbf{x}_{t-2} + \sqrt{1 - \alpha_t\alpha_{t-1}}\bar{\boldsymbol{\epsilon}}_{t-2} \quad \text{merges two Gaussians} \\ &= \dots \\ &= \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\boldsymbol{\epsilon}\end{aligned}$$

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I})$$

- Enable sampling at any time t
- $q(\mathbf{x}_t|\mathbf{x}_{t-1}) = q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)$  ?

# Reverse Diffusion Process

- Compute  $q(x_{t-1}|x_t, x_0)$

$$\begin{aligned} q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) &= q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) \frac{q(\mathbf{x}_{t-1}|\mathbf{x}_0)}{q(\mathbf{x}_t|\mathbf{x}_0)} \\ &\propto \exp \left( -\frac{1}{2} \left( \frac{(\mathbf{x}_t - \sqrt{\alpha_t} \mathbf{x}_{t-1})^2}{\beta_t} + \frac{(\mathbf{x}_{t-1} - \sqrt{\bar{\alpha}_{t-1}} \mathbf{x}_0)^2}{1 - \bar{\alpha}_{t-1}} - \frac{(\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \mathbf{x}_0)^2}{1 - \bar{\alpha}_t} \right) \right) \\ &= \exp \left( -\frac{1}{2} \left( \frac{\mathbf{x}_t^2 - 2\sqrt{\alpha_t} \mathbf{x}_t \mathbf{x}_{t-1} + \alpha_t \mathbf{x}_{t-1}^2}{\beta_t} + \frac{\mathbf{x}_{t-1}^2 - 2\sqrt{\bar{\alpha}_{t-1}} \mathbf{x}_0 \mathbf{x}_{t-1} + \bar{\alpha}_{t-1} \mathbf{x}_0^2}{1 - \bar{\alpha}_{t-1}} - \frac{(\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \mathbf{x}_0)^2}{1 - \bar{\alpha}_t} \right) \right) \\ &= \exp \left( -\frac{1}{2} \left( \left( \frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} \right) \mathbf{x}_{t-1}^2 - \left( \frac{2\sqrt{\alpha_t}}{\beta_t} \mathbf{x}_t + \frac{2\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_0 \right) \mathbf{x}_{t-1} + C(\mathbf{x}_t, \mathbf{x}_0) \right) \right) \end{aligned}$$

- So we have  $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\boldsymbol{\mu}}(\mathbf{x}_t, \mathbf{x}_0), \tilde{\boldsymbol{\beta}}_t \mathbf{I})$

# Reverse Diffusion Process

- Compute  $\tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t$

$$\tilde{\beta}_t = 1 / \left( \frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} \right) = 1 / \left( \frac{\alpha_t - \bar{\alpha}_t + \beta_t}{\beta_t(1 - \bar{\alpha}_{t-1})} \right) = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \cdot \beta_t$$

$$\begin{aligned}\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) &= \left( \frac{\sqrt{\alpha_t}}{\beta_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_0 \right) / \left( \frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} \right) \\ &= \left( \frac{\sqrt{\alpha_t}}{\beta_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_0 \right) \frac{\beta_t(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \\ &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0 \\ &= \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} \frac{1}{\sqrt{\bar{\alpha}_t}} (\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_t) \\ &= \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_t \right)\end{aligned}$$

# ELBO

- Recall the ELBO from VAE:  $\int q(z|x) \log \frac{p_\theta(x, z)}{q(z|x)} dz$
- For the diffusion process:
  - Joint distribution:  $p_\theta(\mathbf{x}_0, \mathbf{x}_{1:T}) = p_\theta(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)\prod_{t=2}^T p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$
  - Posterior:  $q_\phi(\mathbf{x}_{1:T}|\mathbf{x}_0) = q_\phi(\mathbf{x}_1|\mathbf{x}_0)\prod_{t=2}^T q_\phi(\mathbf{x}_t|\mathbf{x}_{t-1})$
  - Substitute to get the ELBO for diffusion:

$$\mathbb{E}_{q_\phi(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_0, \mathbf{x}_{1:T})}{q_\phi(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] = \mathbb{E}_{q_\phi(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)\prod_{t=2}^T p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_\phi(\mathbf{x}_1|\mathbf{x}_0)\prod_{t=2}^T q_\phi(\mathbf{x}_t|\mathbf{x}_{t-1})} \right]$$

# ELBO Derivation

- Compute  $q_\phi(\mathbf{x}_t | \mathbf{x}_{t-1}) = q_\phi(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) = \frac{q_\phi(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)q_\phi(\mathbf{x}_t | \mathbf{x}_0)}{q_\phi(\mathbf{x}_{t-1} | \mathbf{x}_0)}$
- Substitute into the ELBO:

$$\begin{aligned}\log p(\mathbf{x}) &\geq \mathbb{E}_{q_\phi(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_T)p_\theta(\mathbf{x}_0 | \mathbf{x}_1)\prod_{t=2}^T p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)}{q_\phi(\mathbf{x}_1 | \mathbf{x}_0)\prod_{t=2}^T q_\phi(\mathbf{x}_t | \mathbf{x}_{t-1})} \right] \\&= \mathbb{E}_{q_\phi(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_T)p_\theta(\mathbf{x}_0 | \mathbf{x}_1)\prod_{t=2}^T p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)}{q_\phi(\mathbf{x}_1 | \mathbf{x}_0)\prod_{t=2}^T q_\phi(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)q_\phi(\mathbf{x}_t | \mathbf{x}_0)/q_\phi(\mathbf{x}_{t-1} | \mathbf{x}_0)} \right] \\&= \mathbb{E}_{q_\phi(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_T)p_\theta(\mathbf{x}_0 | \mathbf{x}_1)\prod_{t=2}^T p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)q_\phi(\mathbf{x}_1 | \mathbf{x}_0)}{q_\phi(\mathbf{x}_1 | \mathbf{x}_0)\prod_{t=2}^T q_\phi(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)q_\phi(\mathbf{x}_T | \mathbf{x}_0)} \right] \\&= \mathbb{E}_{q_\phi(\mathbf{x}_{1:T} | \mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_T)p_\theta(\mathbf{x}_0 | \mathbf{x}_1)q_\phi(\mathbf{x}_1 | \mathbf{x}_0)}{q_\phi(\mathbf{x}_1 | \mathbf{x}_0)q_\phi(\mathbf{x}_T | \mathbf{x}_0)} \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)}{q_\phi(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)} \right]\end{aligned}$$

# ELBO Derivation

- Reorganize into three terms

$$\begin{aligned}\log p(\mathbf{x}) &\geq \mathbb{E}_{q_\phi(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_T)p_\theta(\mathbf{x}_0|\mathbf{x}_1)}{q_\phi(\mathbf{x}_T|\mathbf{x}_0)} \prod_{t=2}^T \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_\phi(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\&= \mathbb{E}_{q_\phi(\mathbf{x}_{1:T}|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)] + \mathbb{E}_{q_\phi(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_T)}{q_\phi(\mathbf{x}_T|\mathbf{x}_0)} \right] + \sum_{t=2}^T \mathbb{E}_{q_\phi(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_\phi(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\&= \underbrace{\mathbb{E}_{q_\phi(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)]}_{\text{Reconstruction term}} + \mathbb{E}_{q_\phi(\mathbf{x}_T|\mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_T)}{q_\phi(\mathbf{x}_T|\mathbf{x}_0)} \right] + \sum_{t=2}^T \mathbb{E}_{q_\phi(\mathbf{x}_t|\mathbf{x}_0)} \left[ \log \frac{p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q_\phi(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\&= \underbrace{\mathbb{E}_{q_\phi(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)]}_{L_0} - \underbrace{D_{\text{KL}}(q_\phi(\mathbf{x}_T|\mathbf{x}_0) \| p_\theta(\mathbf{x}_T))}_{\text{Prior matching term}} - \underbrace{\sum_{t=2}^T \mathbb{E}_{q_\phi(\mathbf{x}_t|\mathbf{x}_0)} [D_{\text{KL}}(q_\phi(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \| p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))]}_{\text{Score matching term}} \\&= \underbrace{\mathbb{E}_{q_\phi(\mathbf{x}_1|\mathbf{x}_0)} [\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)]}_{L_0} - \underbrace{D_{\text{KL}}(q_\phi(\mathbf{x}_T|\mathbf{x}_0) \| p_\theta(\mathbf{x}_T))}_{L_T} - \sum_{t=2}^T \underbrace{\mathbb{E}_{q_\phi(\mathbf{x}_t|\mathbf{x}_0)} [D_{\text{KL}}(q_\phi(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \| p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))]}_{L_{t-1}}\end{aligned}$$

# ELBO Derivation

- Recall  $D_{\text{KL}}(\mathcal{N}_0 \| \mathcal{N}_1) = \frac{1}{2} \left[ \text{tr}(\Sigma_1^{-1} \Sigma_0) - k + (\mu_1 - \mu_0)^T \Sigma_1^{-1} (\mu_1 - \mu_0) + \ln \left( \frac{\det \Sigma_1}{\det \Sigma_0} \right) \right]$
- The score matching term at timestep t in [2, T]:

$$\begin{aligned} L_{t-1} &= \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[ \frac{1}{2\|\Sigma_\theta(\mathbf{x}_t, t)\|_2^2} \|\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) - \mu_\theta(\mathbf{x}_t, t)\|^2 \right] \\ &= \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[ \frac{1}{2\|\Sigma_\theta\|_2^2} \left\| \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_t \right) - \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) \right\|^2 \right] \\ &= \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[ \frac{(1 - \alpha_t)^2}{2\alpha_t(1 - \bar{\alpha}_t)\|\Sigma_\theta\|_2^2} \|\epsilon_t - \epsilon_\theta(\mathbf{x}_t, t)\|^2 \right] \\ &= \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[ \frac{(1 - \alpha_t)^2}{2\alpha_t(1 - \bar{\alpha}_t)\|\Sigma_\theta\|_2^2} \|\epsilon_t - \epsilon_\theta(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon_t, t)\|^2 \right] \end{aligned}$$

- What about the other two terms?  $L_T, L_0$

# DDPM algorithm

□ In practice  $L_{\text{simple}}(\theta) := \mathbb{E}_{t, \mathbf{x}_0, \epsilon} \left[ \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2 \right]$

---

## Algorithm 1 Training

---

```
1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
       $\nabla_\theta \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2$ 
6: until converged
```

---

## Algorithm 2 Sampling

---

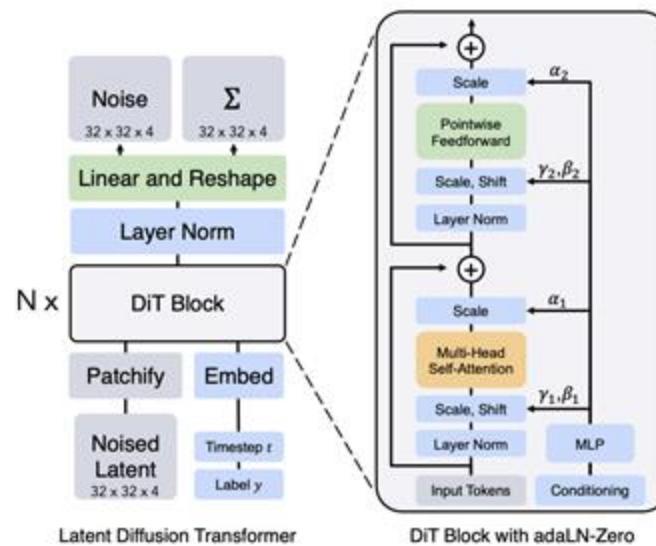
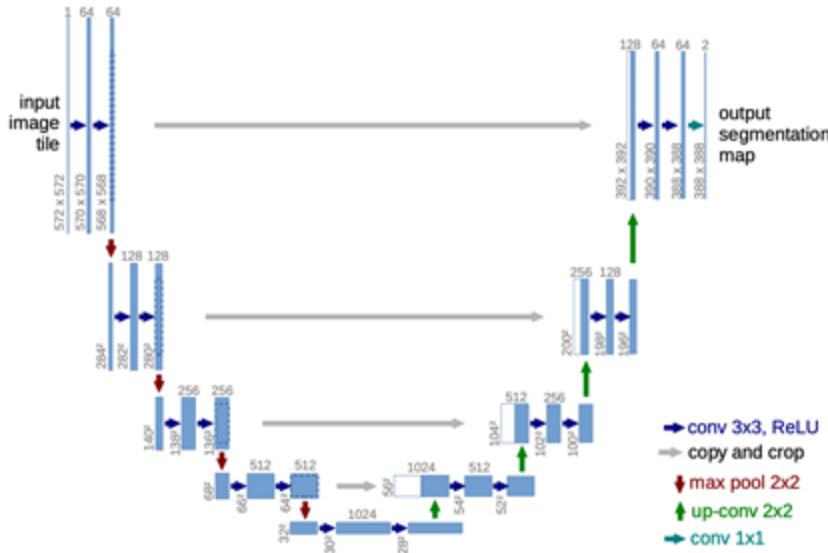
```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\bar{\alpha}_t}} \left( \mathbf{x}_t - \frac{1 - \bar{\alpha}_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

---

- We usually choose T to be a large number, e.g. 1k, 2k, 4k to have better performance
- Sampling is expensive. DDIM, consistency models, distillation...

# Backbones - UNet & DiT

- Conditioning methods: FiLM, AdaLN



# Classifier-Guided Diffusion

- Score function:  $\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t) = \nabla_{\mathbf{x}_t} \log \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}) = -\frac{\mathbf{x}_t - \sqrt{\bar{\alpha}_t} \mathbf{x}_0}{1 - \bar{\alpha}_t} = -\frac{\epsilon_t}{\sqrt{1 - \bar{\alpha}_t}}$

- Joint distribution of data samples and class labels:  $q(\mathbf{x}_t, y)$

- Score function for the joint distribution:

$$\begin{aligned}\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t, y) &= \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t) + \nabla_{\mathbf{x}_t} \log q(y|\mathbf{x}_t) \\ &\approx -\frac{1}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) + \nabla_{\mathbf{x}_t} \log f_\phi(y|\mathbf{x}_t) \quad \text{Trained classifier} \\ &= -\frac{1}{\sqrt{1 - \bar{\alpha}_t}} (\epsilon_\theta(\mathbf{x}_t, t) - \sqrt{1 - \bar{\alpha}_t} \nabla_{\mathbf{x}_t} \log f_\phi(y|\mathbf{x}_t))\end{aligned}$$

- New classifier-guided noise predictor:

$$\bar{\epsilon}_\theta(\mathbf{x}_t, t) = \epsilon_\theta(\mathbf{x}_t, t) - \sqrt{1 - \bar{\alpha}_t} w \nabla_{\mathbf{x}_t} \log f_\phi(y|\mathbf{x}_t) \quad \longleftarrow \quad \text{Ablated diffusion model (ADM)}$$

# Classifier-Free-Guided Diffusion

- ❑ What if there's no trained classifier?
- ❑ Consider the conditional distribution using Bayes rule:

$$\begin{aligned}\nabla_{\mathbf{x}_t} \log p(y|\mathbf{x}_t) &= \nabla_{\mathbf{x}_t} \log p(\mathbf{x}_t|y) - \nabla_{\mathbf{x}_t} \log p(\mathbf{x}_t) \\ &= -\frac{1}{\sqrt{1-\bar{\alpha}_t}} (\epsilon_\theta(\mathbf{x}_t, t, y) - \epsilon_\theta(\mathbf{x}_t, t, y = \emptyset))\end{aligned}$$

- ❑ Then we have the noise predictor with class labels guidance:

$$\begin{aligned}\bar{\epsilon}_\theta(\mathbf{x}_t, t, y) &= \epsilon_\theta(\mathbf{x}_t, t, y) - \sqrt{1-\bar{\alpha}_t} w \nabla_{\mathbf{x}_t} \log p(y|\mathbf{x}_t) \\ &= \epsilon_\theta(\mathbf{x}_t, t, y) + w(\epsilon_\theta(\mathbf{x}_t, t, y) - \epsilon_\theta(\mathbf{x}_t, t)) \\ &= (w+1)\epsilon_\theta(\mathbf{x}_t, t, y) - w\epsilon_\theta(\mathbf{x}_t, t)\end{aligned}$$

# Imitation Learning as Conditional Generation

- Conditional sampling

Learn  $p_\theta(x | c)$  to sample  $x$  given class labels

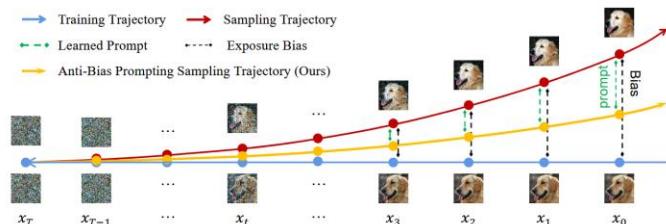
- Maximize the likelihood

$$\max_{\theta} \mathbb{E}_{(x_0, c) \sim \mathcal{D}} [\log p_\theta(x_0 | c)]$$

- Classifier-free guidance

$$p_\theta(x_0 | c)$$

- Exposure bias (diffusion models)



- Conditional sampling

Learn  $\pi_\theta(a | s)$  to sample action given states

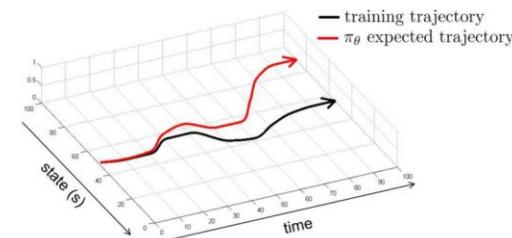
- Maximize the likelihood

$$\max_{\theta} \mathbb{E}_{(s, a) \sim \mathcal{D}} [\log \pi_\theta(a | s)]$$

- Goal-conditioned BC

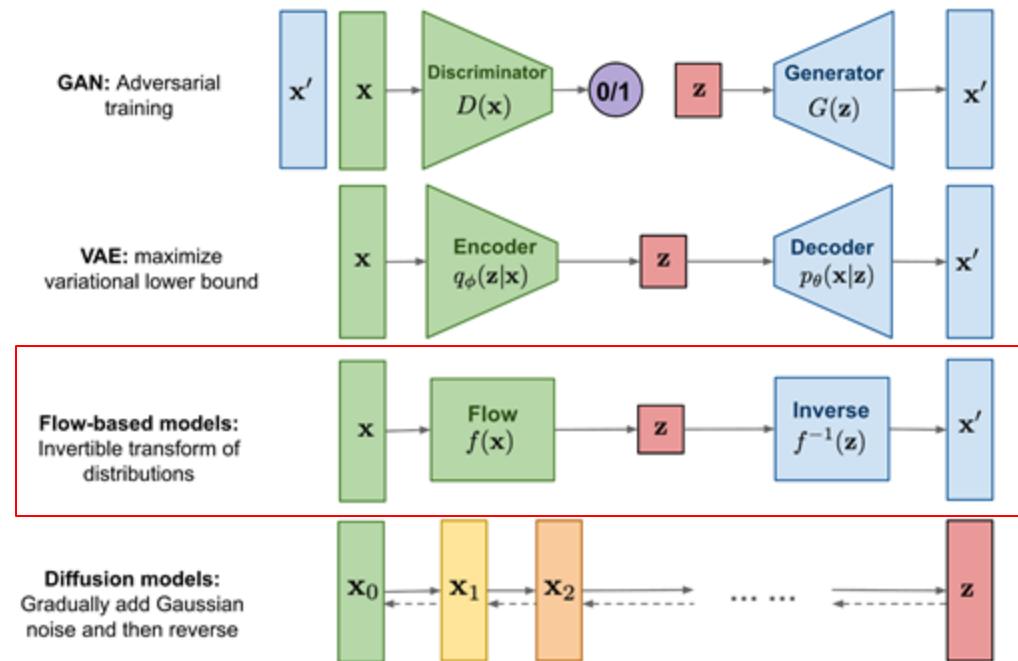
$$\pi_\theta(a | s, g)$$

- Distribution shift

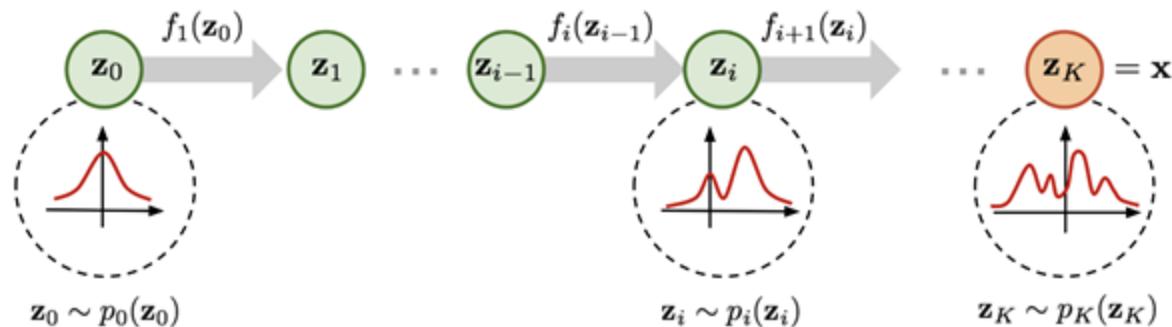


# Generative Models

- ❑ Variational Autoencoder
- ❑ Diffusion Models
- ❑ Flow-based Models



# Normalizing Flow



- Definition:  $\mathbf{z}_{i-1} \sim \pi(\mathbf{z}_{i-1}), \mathbf{z}_i = f_i(\mathbf{z}_{i-1}), \mathbf{z}_{i-1} = f_i^{-1}(\mathbf{z}_i)$
- Compute the data distribution:  $\mathbf{x} = \mathbf{z}_K = f_K \circ f_{K-1} \circ \dots \circ f_1(\mathbf{z}_0)$
- Learn by maximizing the log-likelihood
- How to compute the log-likelihood?  $\log p(\mathbf{x}) = \log \pi_K(\mathbf{z}_K)$

# Normalizing Flow

- Preliminary: given  $z \sim \pi(z)$ , construct new variable  $x = f(z)$ ,  $f$  is invertible, then we have

$$\int p(x)dx = \int \pi(z)dz = 1 \quad p(x) = \pi(z) \left| \frac{dz}{dx} \right| = \pi(f^{-1}(x)) \left| \frac{df^{-1}}{dx} \right| = \pi(f^{-1}(x)) |(f^{-1})'(x)|$$

- For multivariate case:  $\mathbf{z} \sim \pi(\mathbf{z}), \mathbf{x} = f(\mathbf{z}), \mathbf{z} = f^{-1}(\mathbf{x})$

$$p(\mathbf{x}) = \pi(\mathbf{z}) \left| \det \frac{d\mathbf{z}}{d\mathbf{x}} \right| = \pi(f^{-1}(\mathbf{x})) \left| \det \frac{df^{-1}}{d\mathbf{x}} \right|$$

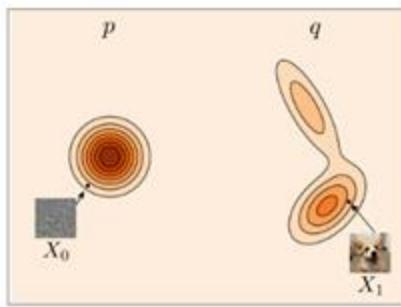
- Inverse function theorem: given  $y = f(x)$  and  $x = f^{-1}(y)$ , then we have

$$\frac{df^{-1}(y)}{dy} = \frac{dx}{dy} = \left( \frac{dy}{dx} \right)^{-1} = \left( \frac{df(x)}{dx} \right)^{-1}$$

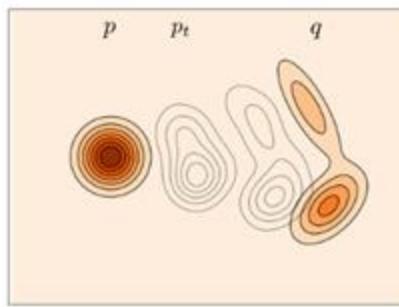
# Normalizing Flow

- Given  $\mathbf{z}_{i-1} \sim p_{i-1}(\mathbf{z}_{i-1})$ ,  $\mathbf{z}_i = f_i(\mathbf{z}_{i-1})$ ,  $\mathbf{z}_{i-1} = f_i^{-1}(\mathbf{z}_i)$
- Compute  $\log p(\mathbf{x}) = \log \pi_K(\mathbf{z}_K)$
- From the change of variable theorem:  $p_i(\mathbf{z}_i) = p_{i-1}(f_i^{-1}(\mathbf{z}_i)) \left| \det \left( \frac{df_i^{-1}}{d\mathbf{z}_i} \right) \right|$ 
$$= p_{i-1}(\mathbf{z}_{i-1}) \left| \det \left( \left( \frac{df_i}{d\mathbf{z}_{i-1}} \right)^{-1} \right) \right| = p_{i-1}(\mathbf{z}_{i-1}) \left| \left( \det \left( \frac{df_i}{d\mathbf{z}_{i-1}} \right) \right)^{-1} \right| = p_{i-1}(\mathbf{z}_{i-1}) \frac{1}{\left| \det \left( \frac{df_i}{d\mathbf{z}_{i-1}} \right) \right|}$$
- We get  $\log p_i(\mathbf{z}_i) = \log p_{i-1}(\mathbf{z}_{i-1}) - \log \left| \det \left( \frac{df_i}{d\mathbf{z}_{i-1}} \right) \right|$
- And  $\log p(\mathbf{x}) = \log \pi_K(\mathbf{z}_K)$ 
$$= \log \pi_{K-1}(\mathbf{z}_{K-1}) - \log \left| \det \frac{df_K}{d\mathbf{z}_{K-1}} \right|$$
$$= \log \pi_{K-2}(\mathbf{z}_{K-2}) - \log \left| \det \frac{df_{K-1}}{d\mathbf{z}_{K-2}} \right| - \log \left| \det \frac{df_K}{d\mathbf{z}_{K-1}} \right| = \log \pi_0(\mathbf{z}_0) - \sum_{i=1}^K \log \left| \det \frac{df_i}{d\mathbf{z}_{i-1}} \right|$$
- Requirements on  $f$ ?
  - 1. Easily invertible. 2. Easy to compute jacobians

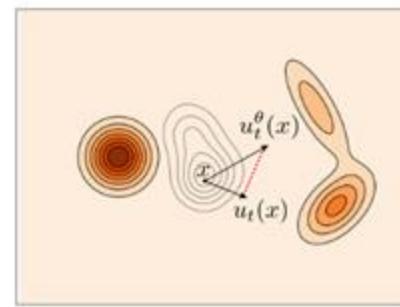
# Flow-Matching Methods



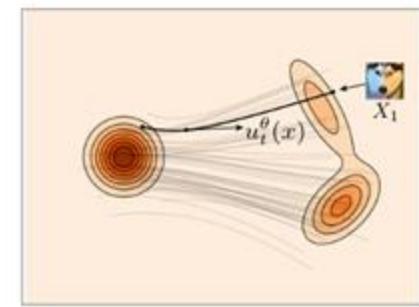
(a) Data.



(b) Path design.



(c) Training.



(d) Sampling.

- Training:
  - Build a probability path  $(p_t)_{0 \leq t \leq 1}$  from a **known** source distribution  $p$  to a **target** distribution  $q$
  - regression on the **vector field** used to convert distributions along the prob path
- Sampling (from the target distribution):
  - Sample from the source distribution  $X_0 \sim p$
  - Solve an ODE determined by the vector field to get  $X_1 \sim q$

# Flow-Matching Methods

- Vector field:  $u : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$ , flow:  $\psi : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$
- ODE:  $\frac{d}{dt}\psi_t(x) = u_t(\psi_t(x))$  where  $\psi_t := \psi(t, x)$  and  $\psi_0(x) = x$
- $u$  generates the prob. path if  $X_t := \psi_t(X_0) \sim p_t$  for  $X_0 \sim p_0$
- **Learning objective:** learn a vector field that can generates the prob. path  $p_t$
- **A simple probability path?**  $X_t = tX_1 + (1 - t)X_0 \sim p_t$
- **Flow matching loss:**  $\mathcal{L}_{\text{FM}}(\theta) = \mathbb{E}_{t, X_t} \|u_t^\theta(X_t) - u_t(X_t)\|^2$ , where  $t \sim \mathcal{U}[0, 1]$  and  $X_t \sim p_t$
- **What's the issue?**
  - We cannot compute the target distribution  $p_1$

# Flow-Matching Methods

- Conditional random variables:  $X_{t|1} = tx_1 + (1 - t)X_0 \sim p_{t|1}(\cdot | x_1) = \mathcal{N}(\cdot | tx_1, (1 - t)^2 I)$
- Solving for the cond. vector field:  $\frac{d}{dt}X_{t|1} = u_t(X_{t|1}|x_1) \longrightarrow u_t(x|x_1) = \frac{x_1 - x}{1 - t}$
- **Conditional flow matching loss:**

$$\mathcal{L}_{\text{CFM}}(\theta) = \mathbb{E}_{t, X_t, X_1} \|u_t^\theta(X_t) - u_t(X_t|X_1)\|^2, \text{ where } t \sim U[0, 1], X_0 \sim p, X_1 \sim q$$



$$\mathcal{L}_{\text{CFM}}^{\text{OT, Gauss}}(\theta) = \mathbb{E}_{t, X_0, X_1} \|u_t^\theta(X_t) - (X_1 - X_0)\|^2, \text{ where } t \sim U[0, 1], X_0 \sim \mathcal{N}(0, I), X_1 \sim q$$

# References

- Weng, Lilian. (Jul 2021). What are diffusion models? Lil'Log. <https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>.
- Lipman, Yaron, et al. "Flow matching guide and code." arXiv preprint arXiv:2412.06264 (2024).
- Jonathan Ho et al. "[Denoising diffusion probabilistic models.](#)" arxiv Preprint arxiv:2006.11239 (2020).